



Multi-Layer Flake Architecture for Holochain + Dynamo Systems

Integrating a Holochain **user-layer** (peer-to-peer app conductors) with a FlexNetOS/Dynamo **organization-layer** (orchestration and infrastructure services) can be achieved with NixOS flakes using a modular, multi-component design. Recent examples in the Holochain community illustrate how to structure such a system with multiple flake modules, containerized services, and clear integration boundaries between layers.

Modular Flake Design with Multiple Layers

A good starting point is a flake that contains separate NixOS modules or configurations for each “layer” or component. Tools like **flake-parts** (or Numtide’s Blueprint) allow composing many modules into one flake output. For example, the `kickstart.nix` template uses flake-parts to aggregate numerous preset modules (for different languages, systems, etc.) under one flake [1](#) [2](#). Similarly, the Holochain repository itself employs flake-parts to auto-import all Nix modules in a directory, merging them into the flake outputs [3](#). This modular approach means you can define separate NixOS configurations for the Holochain conductor node(s) and for the Dynamo orchestration services, then compose them in the flake.

In practice, you might have one NixOS module for the **Holochain conductor** (the user-layer, potentially running as a service or in a container) and another for the **Dynamo orchestration node** (the organization-layer logic). The flake can expose multiple `nixosConfigurations` or profiles. For instance, a monorepo approach used by the Holo-Host project defines multiple system roles in one flake and uses a Nix module framework to include the needed pieces for each [3](#). Each role (e.g. conductor, orchestrator, gateway) can be a flake module that is enabled or disabled depending on the target system, or even different NixOS host definitions in `nixosConfigurations`. This keeps layer-specific configuration isolated but still allows a unified build/deploy process.

Real-World Example: Holochain Hosting (Holo-Host)

One of the most relevant examples is **Holo-Host**, which combines a Holochain conductor with an orchestration layer in a Nix-based system. Holo-Host’s architecture (circa 2023–2025) is a monorepo with a flake that defines multiple services and uses NixOS containers for isolation. In their development environment, they spin up four logical components as separate systemd services/containers [4](#):

- `dev-hub` – a NATS messaging server (also acts as a bootstrap for peer discovery) [5](#)
- `dev-orch` – the orchestrator service (organization-layer logic) [5](#)
- `dev-host` – the Holochain host agent, which runs a Holochain conductor for hApps (user-layer) [5](#)
- `dev-gw` – a gateway service (providing an external API/HTTP interface to the system) [5](#)

Each of these runs in its own lightweight NixOS container (via `systemd-nspawn`) or process, defined by NixOS modules. The flake uses a tool (`blueprint`) to manage these modules, and starting the system

brings up all components. A single command (via `just`) launches all dev containers, sets up the NATS infrastructure, and initializes the Holochain conductor ⁶. This demonstrates a pattern where the flake coordinates multiple NixOS configurations at once, providing a full multi-service environment.

Service isolation is achieved by running the Holochain conductor in a separate container (the host-agent container) with its own filesystem and network namespace. The orchestrator and other org-layer services run in other containers or on the host. This separation is important for security (the Holochain conductor has its own sandbox) and manageability. They communicate through well-defined channels rather than direct function calls.

Inter-Layer Messaging and Integration Boundaries

In a Holochain+Dynamo stack, the layers typically communicate via networking or messaging queues, not direct in-process calls. Holo-Host uses **NATS** as a lightweight message bus between the organization layer and the Holochain host agents ⁷ ⁸. For example, when the orchestrator needs to install or update a hApp on a host, it publishes commands over NATS, which the host agent (listening on NATS) receives and acts on ⁹. This decouples the two layers – the orchestrator doesn't invoke Holochain internals directly; instead it sends messages (e.g. "install this DNA") and the host agent handles the Holochain conductor API. This messaging approach is a common integration boundary: it keeps the user-layer logic (running on possibly many distributed nodes) separate from org-layer coordination logic, with a clear protocol between them (e.g. NATS messages or gRPC calls). Other systems might use Apache Kafka or REST, but NATS is noted for simplicity in this context ¹⁰ ¹¹.

Another integration boundary is **networking and service APIs**. The Holochain conductor typically exposes a local admin interface (IPC or localhost-bound HTTP/WebSocket for control) and perhaps a public interface for hApp web UIs. In Holo-Host, since the conductor runs inside an isolated container and binds only to localhost, they implemented a clever **two-tier port forwarding** with `socat` to bridge the gap ¹². Inside each Holochain container, an internal socat process forwards the conductor's localhost-only port to a container-wide accessible port. Then on the host, another socat forwards from the host machine into the container's forwarded port ¹³. This allows the orchestration layer or a gateway to talk to the Holochain admin interface securely, without exposing that interface broadly. It's an example of handling the integration boundary between isolated environments: Holochain stays locked to localhost inside its sandbox, and only a controlled tunnel provides access for the outside services ¹⁴ ¹⁵. In production, NixOS's container module can set this up (as shown by enabling `containerPrivateNetwork = true` and having the service generate the socat tunnels automatically ¹⁶).

Service interfaces are another boundary consideration. The org-layer (Dynamo) might expose an API (HTTP/gRPC as described in the NOA design) and the user-layer (Holochain) exposes zome calls via an HTTP gateway or directly to UIs. In Holo-Host, a dedicated **gateway service** (`dev-gw`) acts as an HTTP bridge to the Holochain apps ¹⁷, so that external clients can hit a REST endpoint which the gateway translates into Holochain function calls (often via the conductor's HTTP interface or through the host agent). This separation means each layer has its own responsibilities and they only meet at specific API points (e.g. messaging bus, HTTP interface, etc.). The **FlexNetOS/Dynamo** layer would similarly interact with Holochain either by running its own conductor instances or sending requests to conductors on user devices. Keeping those interactions limited to network protocols (like NATS messages or HTTP calls) enforces a clean integration boundary.

Runtime Orchestration: Systemd, Containers, Remote Builds

For deploying and running such a composite system, NixOS's strengths come into play:

- **Systemd services:** You can package the Holochain conductor as a NixOS module that runs `holochain` (or your DNA) as a systemd service, and package Dynamo's components as their own services. Systemd provides supervision, logging, and isolation (via cgroups, user permissions, etc.). This is simpler if everything runs on one host. If using multiple machines, each machine's NixOS config (from the flake) would include the relevant services for that role.
- **NixOS Containers:** As seen in Holo-Host, NixOS's built-in container support (`systemd-nspawn`) allows you to declaratively define containers in your NixOS config (`containers.<name>.config = { ... }`). Each container runs a nested NixOS system. This is great for simulating a multi-machine environment on one host or enforcing stronger isolation. It appears Holo-Host uses this for dev/testing (and possibly in production on HoloPort devices)¹⁸ ¹⁹. You could run the Holochain user layer in a container on an organization-controlled server for security, for example. Another alternative is lightweight VMs (with firecracker or QEMU via `nixosTests`) if stronger isolation is needed; NixOS tests often use such VMs to run multi-node scenarios.
- **Remote builders and deployers:** While not specific to runtime orchestration, using Nix's remote build capabilities can improve how you build and deploy this layered system. For instance, you might use a **remote builder** machine or CI (like Hydra or Buildbot) to build the whole flake, producing system closures for each layer. Holo-Host's CI uses `buildbot-nix` to coordinate builds²⁰ – this ensures that all services (Rust binaries, etc.) are built in a consistent environment and cached. Then, deployment tools like `deploy-rs` or `Colmena` can push the NixOS configurations to the target machines or activate them. Colmena, for example, allows describing a fleet of NixOS hosts in one flake and deploying in one command, which is useful if your "org layer" runs on one server and the "user layer" on edge devices or containers. By late 2025, these tools are well-established for Nix-based devops.

In a scenario where the user-layer might be on many devices (like user-run agents) that you don't directly control with NixOps, an organization might instead distribute Holochain apps to users via other means (app store, etc.) and only orchestrate certain **organization-hosted Holochain instances**. If FlexNetOS Dynamo is coordinating a fleet of servers (perhaps similar to Holo's cloud hosting), using Nix-based remote deployment ensures that the same flake can define both the base OS image and the Holochain services on those hosts.

Best Practices and Evolving Patterns (2025+)

1. **Use Flakes to enforce consistency across layers:** Both layers can share inputs (e.g. the same `nixpkgs` pin or Rust toolchains). For example, Holochain's flake pins a specific Holonix version (which bundles Holochain, lair, etc.) so that the conductor and hApp build environment are in sync²¹. Your flake can similarly include specific versions of Holochain binaries and Dynamo's dependencies to avoid drift between layers.

2. Modularize configurations: Structure your flake with separate modules for each major service. This improves clarity and allows testing or deploying components independently. Holochain's flake and Kickstart.nix both illustrate how to split configuration for manageability ³ ². You can have, for instance, a `holochain-conductor.nix` module enabling the conductor service (with options for which DNA to run, ports, etc.), and a `dynamo-org.nix` module for the org-layer daemon (with its database config, message broker settings, etc.). These can then be composed in a top-level `configuration.nix` depending on the role of the machine.

3. Leverage messaging and APIs, not tight coupling: Define a clear API between Holochain and Dynamo. The Holo-Host approach of using NATS and capability tokens (or HTTP + JWT) is instructive – it treats the Holochain node almost like an external microservice. In your design, you might use an **async message queue** or **RPC calls** to cross the boundary. The FlexNetOS "NOA" spec indicates preference for NATS for simplicity ¹⁰ ¹¹, which aligns well with what Holochain hosts have done. This decoupling allows each layer to evolve or restart independently (for example, you can update the orchestrator without restarting all Holochain nodes, as long as they reconnect to NATS).

4. Use systemd and cgroups for resource isolation: If running multiple components on one host (or one NixOS system), consider using systemd's features (service `PrivateTmp`, `MemoryMax`, `CPUQuota`, etc.) to isolate resources between Holochain and Dynamo processes. This is especially important if the user-layer can consume unpredictable resources (e.g. heavy DHT workloads or WASM execution). Running Holochain in a container or VM is the strongest isolation (as Holo does), but even within one OS, systemd knobs can prevent one service from starving others.

5. Consider containerization vs. orchestration tools: As of 2025, many NixOS deployments for complex systems either use built-in containers (as Holo does) or integrate with container runtimes/Kubernetes. There aren't many public examples of Holochain on Kubernetes yet, but it's conceivable. If FlexNetOS Dynamo has Kubernetes in its scope, you could still use Nix to produce container images for Holochain and Dynamo services. However, the simplicity of NixOS containers and systemd, as used in Holo-Host, has proven effective for a distributed P2P system. It avoids the overhead of a full Kubernetes stack and keeps everything declarative in one flake.

6. New community tools: The Nix community has been moving toward higher-level frameworks for config. Projects like **DevOS** and **Morph** (and others like **Nixos-General** templates) provide patterns for multi-host flakes. While not Holochain-specific, they can help structure a flake that targets multiple machines/roles easily. Additionally, **impermanence.nix** is a module that can be useful for Holochain, as it lets you declare certain directories (like the conductor's data storage) as ephemeral or on tmpfs, which might be desirable for stateless deployments or testing.

In summary, the **state of the art** for layering a Holochain user network with an org-level compute layer using Nix flakes is exemplified by the Holochain+Holo approach. A unified flake can contain all the pieces: you get reproducible builds of both the Holochain binaries and the Dynamo services, and you can spin up the entire stack consistently. By following patterns of module separation, containerized isolation via `systemd-nspawn`, and message-based integration (NATS/HTTP), you achieve a robust separation of concerns. This prevents the user-layer from compromising the org-layer while still allowing them to work in concert through well-defined channels. As seen with Holo-Host's containers and two-tier forwarding setup, these practices handle service isolation and networking quirks gracefully (ensuring, for example, that Holochain's local interfaces are reachable only by the intended orchestrator) ¹³ ²².

Moving forward, adopting such **flake-based orchestration** means your entire distributed stack is defined as code. This not only improves reliability (no configuration drift between nodes) but also makes it easier to test complex scenarios (using NixOS VM tests or containers to simulate a network). The community is increasingly sharing NixOS modules for services (e.g. a community NixOS module for running Holochain could emerge, similar to how there are modules for IPFS, etc.), which would further simplify integration. As of late 2025, using flakes with multi-module structures and NixOS containers is a proven, recommended approach for combining peer-to-peer agents with centralized orchestration in a single, coherent system configuration.

Sources

- ALT-F4-LLC's **Kickstart Nix** flake – demonstrates a modular flake template using flake-parts to combine multiple language, system, and Home Manager modules ① ② .
- **Holochain core** using flake-parts – auto-imports all Nix modules for different concerns, illustrating how to merge many configs in one flake ③ .
- **Holo-Host (Holochain hosting)** monorepo – a real-world multi-service NixOS flake. Defines orchestrator, NATS, host agent, and gateway as separate modules/containers ④ . Uses systemd-spawn containers with private networking and socat tunnels to safely expose Holochain's interfaces ⑯ . Employs NATS messaging for inter-service communication and workload commands ⑬ ⑭ . This architecture shows how to layer a user-centric Holochain conductor within an org-managed environment using Nix.
- FlexNetOS NOA Ark documentation – indicates the use of NATS for messaging and a Rust-first, modular approach to orchestration, aligning with the above patterns ⑩ ⑪ .

① ② [flake.nix](#)

<https://github.com/ALT-F4-LLC/kickstart.nix/blob/96fb4e38571faf93bce1fd14fd738a84dbbdac1d/flake.nix>

③ [flake.nix](#)

<https://github.com/holochain/holochain/blob/926d6340e00bcd7123b756ee1c4d3754dc575b56/flake.nix>

④ ⑤ ⑥ ⑦ ⑧ ⑨ ⑫ ⑬ ⑭ ⑮ ⑯ ⑰ ⑱ ⑲ ⑳ ⑳ [GitHub - Holo-Host/holo-host: Holo Host main repository](#)

<https://github.com/Holo-Host/holo-host>

⑩ ⑪ [BUILD_SPEC.md](#)

https://github.com/FlexNetOS/noa_ark_os/blob/2c00cb0cf28d4ade476544e4a951cc56df848d27/server/BUILD_SPEC.md